

How should we **measure biases & societal phenomena** from text?

Like in what extent *nurse* is perceived as a *women job* when processing Wikipedia text

- Widely accepted approach to bias measurement is based on **high-order relations** (like cosine)
- We suggest **smoothed first-order measurement**
 - achieved by reconstructing the co-occurrence matrices from word2vec/GloVe word embeddings
- Smoothed first order bias measurement has several **advantages**:
 - Avoid the inclusion of **unwanted factors** as in high-order methods
 - Approaches the **sparsity issue** of traditional count-based methods
- Our experiments reveal the existence of **stronger biases** in occupations **towards female** when processing Wikipedia text
- A Take-Home Message: select a *proper bias measurement method* based on what is aimed to be quantified as *bias construct*

